

Yicun Yang

Ph.D. Student in Artificial Intelligence · HKUST(GZ)
yangyicun187@gmail.com · maomaocun.github.io

Education

- **Ph.D. in Artificial Intelligence**, The Hong Kong University of Science and Technology (Guangzhou), AI Thrust, 2026–present. Primarily supervised by Prof. Linfeng Zhang and co-supervised by Prof. Xuming Hu.
- **B.Eng. in Software Engineering**, Harbin Institute of Technology, 2022–2026.

Research experience

- **Research Intern**, Efficient and Precision Intelligent Computing (EPIC) Lab, Shanghai Jiao Tong University, 2024–2026. Supervised by Prof. Linfeng Zhang.

Research interests

Discrete diffusion language models · efficient LLM inference · adaptive caching · block-wise denoising · speculative decoding · efficient training and serving.

Open source

- [dLLM-Cache](#) — adaptive caching for diffusion LLM inference.
- [dLLM-Var](#) — native variable-length generation for diffusion LLMs.
- [dLLM-Factory](#) — tools for training diffusion language models.

Publications

- [LLaDA2.2: Enabling Agentic Diffusion Language Models via Levenshtein Editing](#) — Technical report, 2026
My name appears as “Yichun Yang” in the technical report due to a typo.
- [dLLM-Cache: Accelerating Diffusion Large Language Models with Adaptive Caching](#) — ICML, 2026
- [Diffusion LLM with Native Variable Generation Lengths: Let \[EOS\] Lead the Way](#) — arXiv, 2025
- [Mask Tokens as Prophet: Fine-Grained Cache Eviction for Efficient dLLM Inference](#) — arXiv, 2025
- [dVLA: Diffusion Vision-Language-Action Model with Multimodal Chain-of-Thought](#) — arXiv, 2025
- [The Devil behind the Mask: An Emergent Safety Vulnerability of Diffusion LLMs](#) — arXiv, 2025
- [Dataset Distillation with Neural Characteristic Function: A Minmax Perspective](#) — CVPR, 2025

Links

[Google Scholar](#) · [alphaXiv](#) · [GitHub](#)